



CLASS 1: ADVANCED ACTIVATION FUNCTIONS

PART 4

Activation functions play a crucial role in neural networks by introducing non-linear properties that allow the model to learn complex patterns in the data. This session will explore additional advanced activation functions beyond the commonly used Sigmoid, Tanh, and ReLU.



Advanced Activation Functions

These functions are essential for deeper networks that need to model more complex relationships within large or highly varied datasets.

Overview of Additional Functions:

Leaky ReLU:

Improves upon the ReLU function by allowing a small, non-zero gradient when the unit is not active and preventing the dying ReLU problem.

Parametric ReLU (PReLU):

This variation of Leaky ReLU has a coefficient for the negative part of the function that is learned during training, potentially improving learning capability.

Exponential Linear Unit (ELU):

ELU speeds up learning and achieves higher classification accuracy by reducing the vanishing gradient problem with an exponential curve for negative values, unlike ReLU.

Softplus:

A smooth approximation to the ReLU function, defined as $\log(1 + e^x)$, which creates a smooth curve and reduces the likelihood of dead neurons.

Swish:

A self-gated activation function that has been found to work better than ReLU in some deep learning models. It is defined as $x * \text{sigmoid}(\beta x)$.



Importance and Selection

Choosing the right activation function can significantly influence the performance of a neural network. Each function has its strengths and weaknesses, and the choice often depends on the specific requirements of the application and the nature of the problem being solved.



Conclusion

This overview of advanced activation functions introduces more tools available to neural network designers. Understanding these options expands the toolkit for designing networks that can learn from complex, high-dimensional datasets.